

Photoacoustic Reconstruction MVP Report

PLANNED to MVP COMPLETE — Implementation, Verification, and Results

Adebanji Adelowo

4 September 2026

Implementation status

Every stage in IMPLEMENTATION_PLAN.md up through the MVP is real, executed, and verified — not scaffolded. `src/phantoms.py`, `forward_model.py`, `baselines.py`, `reconstruction_net.py`, `evaluate.py` all contain working implementations (no `NotImplementedError` remaining). The full pipeline runs end-to-end: **phantom** → **forward simulation** → **sparse measurements** → **time-reversal** → **learned reconstruction** → **PSNR/SSIM evaluation**.

Verification evidence

- **Stage 1:** dedicated conda env (`photoacoustic`, Python 3.11), `jwave/jax/torch/` `scikit-image` installed; API directly introspected via `inspect` (not trusted from memory or the earlier web search) — this surfaced real facts the original architecture got wrong or left unspecified (PML boundary, not periodic; `simulate_wave_propagation`’s actual return shape). Smoke test passed.
- **Stage 2:** phantom generator — 5/5 tests pass, visually inspected. Design corrected from “Shepp-Logan” to seeded Gaussian blobs (physically appropriate reasoning documented).
- **Stage 3:** forward model — 8/8 tests pass. One real test failure diagnosed (sensor-array radius too small, sensor sat in phantom’s Gaussian tail) and fixed, not hidden. Independent physical validation: the sensor trace shows a textbook bipolar photoacoustic pulse, which emerged from the physics rather than being designed in.
- **Stage 4:** time-reversal baseline — 12/12 tests pass. One real API misunderstanding (`on_grid` returns the full trajectory, not just the final step) diagnosed and fixed. 0.929 correlation with ground truth on a real reconstruction.
- **Stage 5:** training-data pipeline — 56 examples generated in 12.5s, zero seed overlap across splits (verified programmatically, not just by convention).
- **Stage 6:** U-Net — overfit sanity check passed (loss dropped 99.8% on a fixed batch) *before* the full run, per the required gate. Full training: val loss 0.0118 → 0.0005 over 60 epochs (5.2s).
- **Stage 7:** evaluation — 15/15 tests pass overall. Real metrics computed on the held-out test set.

Results (real, measured)

sparsity	TR PSNR	TR SSIM	Learned PSNR	Learned SSIM
16 sensors	18.97 dB	0.706	30.73 dB	0.500
64 sensors	22.43 dB	0.735	33.08 dB	0.517

Failures/limitations — reported honestly, not hidden

PSNR and SSIM disagree on which method wins. The learned model wins decisively on PSNR and, unambiguously, visually (streak artifacts almost entirely removed — `report/mvp_comparison.png`).

It *loses* on whole-image SSIM. This was investigated rather than picking the favorable metric: within the true structure, learned SSIM = 0.936 vs. 0.188 for time-reversal (confirms the visual win); in the background (>85% of image area), the U-Net leaves a small residual “haze” (mean 0.0099 vs. near-zero for both ground truth and time-reversal) that SSIM’s local-variance sensitivity penalizes heavily, inverting the whole-image average. This is a real, diagnosed, known failure mode of MSE-trained restoration networks — not fixed in this pass (that would be strong-version scope), documented in `report/mvp_results.txt` and `report/mvp_ssim_diagnosis.png`.

Two minor process issues, both disclosed and fixed forward-only (no history rewrite): a `.gitignore` gap let `data/*.npz` get committed once before being caught and untracked; a test’s sensor radius needed correcting for the same underlying geometric reason as the Stage 3 fix.

Repository changes

25 files changed, +1196/−138 lines across 9 commits (7270cd6..7f37d50) — full stage-by-stage detail in `IMPLEMENTATION_LOG.md`. New: `IMPLEMENTATION_LOG.md`, `scripts/{smoke_test,generate_training_data,train,evaluate,tests/test_{forward_model,baselines,evaluate}.py`, `report/{dev_*,mvp_*}` figures and results. Modified: `ARCHITECTURE.md` (2 corrections found during implementation), `README.md`, `requirements.txt` (pinned), all `src/*.py` (from stubs to real implementations).

MVP verdict

MVP COMPLETE.

Portfolio verdict

Genuinely strong evidence now exists — but it has not been added to any CV or letter, per instruction. Assessment: the real, honest result (dramatic structural improvement, a diagnosed rather than hidden failure mode, working code with tests) is legitimate CV-worthy evidence for CBI and Heckel specifically (Direct/Strong AVS positions), describable accurately as e.g. *“implemented and evaluated a learned photoacoustic reconstruction pipeline against a time-reversal baseline; investigated a PSNR/SSIM metric disagreement to a diagnosed root cause.”* Recommendation against adding it to Quaini or DSBI yet (their AVS for this project was always Moderate/Weak) or to the long-term Selected Projects portfolio until the strong version exists. This is a recommendation, not an action — no application document was changed.

Next recommended action

Given the six not-yet-sent TUM/Helmholtz applications are still sitting in draft: decide whether to (a) add this evidence to the CBI/Heckel CVs now and send, or (b) pursue the strong version (Tikhonov baseline, background-haze fix) first. Leaning toward (a) — the result is real and honestly reportable — but that is a decision for Adebajji, not one made automatically here.

Prepared 4 September 2026, documenting the implementation pass that moved ~/Documents/GitHub/photoacoustic-reconst from PLANNED to MVP COMPLETE. See `IMPLEMENTATION_LOG.md` in that repository for the full stage-by-stage verification record.